
Variance-Weighted Centroid Selection for Prototype-Based 1-Nearest Neighbor Classification

Anonymous Authors¹

Abstract

We address the problem of selecting a small set of representative prototypes from a large training set for 1-nearest neighbor (1-NN) classification. We propose a variance-weighted centroid selection algorithm that allocates prototypes proportionally based on within-class variance and uses K-Means clustering to identify representative samples. Experiments on MNIST demonstrate that our method significantly outperforms random selection, with advantages becoming more pronounced at higher compression ratios. At $60\times$ compression (1,000 prototypes from 60,000 training samples), our method achieves 92.47% accuracy compared to 88.56% for random selection (+3.91%). At extreme compression ($6,000\times$, only 10 prototypes), our method maintains 67.01% accuracy while random selection degrades to 39.37% (+27.64%).

1. Introduction and Key Idea

The k -nearest neighbor (k -NN) algorithm is a fundamental non-parametric classifier that predicts the label of a test sample based on the labels of its k closest training samples (Cover & Hart, 1967). **Prototype selection** addresses the computational limitation of k -NN by selecting a representative subset of training samples to use for classification.

1.1. Key Idea (High-Level Description)

Our prototype selection method is based on two key observations:

1. **Good prototypes are typical samples, not boundary samples.** While decision boundaries are defined by samples near class interfaces, the best prototypes are

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

cluster centroids that represent the “typical” appearance of each class. Boundary points are often noisy or ambiguous.

2. **Classes with higher variance need more prototypes.** Different classes have different amounts of within-class variation. For example, the digit “1” has consistent appearance, while “2” exhibits more variation. We allocate prototypes proportionally to class variance.

Our algorithm: (1) compute within-class variance for each class, (2) allocate prototype budget M proportionally based on variance, (3) use K-Means clustering within each class to find representative samples.

2. Problem Formulation

Given a training set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$ where $\mathbf{x}_i \in \mathbb{R}^d$ and $y_i \in \{0, 1, \dots, C-1\}$, the prototype selection problem is to find a subset $\mathcal{P} \subset \mathcal{D}$ with $|\mathcal{P}| = M \ll N$ that maximizes classification accuracy on a held-out test set when using 1-NN with \mathcal{P} as the reference set.

The 1-NN classifier predicts:

$$\hat{y} = y_{j^*}, \quad \text{where } j^* = \arg \min_{(\mathbf{x}_j, y_j) \in \mathcal{P}} \|\mathbf{x} - \mathbf{x}_j\|_2 \quad (1)$$

The compression ratio is defined as N/M . For MNIST with $N = 60,000$ training samples, selecting $M = 1,000$ prototypes yields $60\times$ compression.

3. Algorithm

Implementation Details:

- We use MiniBatchKMeans (sklearn) with batch_size=256, n_init=3, max_iter=100 for efficiency.
- Time complexity: $O(N \cdot M \cdot t)$ where t is K-Means iterations.
- Runtime: 2–680 seconds depending on M (on a standard laptop).

Algorithm 1 Variance-Weighted Centroid Selection

Input: Training data $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$, budget M

Output: Prototype indices \mathcal{I}

// Step 1: Compute within-class variance

for each class $c = 0, \dots, C - 1$ **do**

$$\boldsymbol{\mu}_c \leftarrow \frac{1}{|D_c|} \sum_{\mathbf{x} \in D_c} \mathbf{x} \quad \text{(class mean)}$$

$$\sigma_c^2 \leftarrow \frac{1}{|D_c|} \sum_{\mathbf{x} \in D_c} \|\mathbf{x} - \boldsymbol{\mu}_c\|^2 \quad \text{(class variance)}$$

end for

// Step 2: Allocate prototypes proportionally to variance

for each class c **do**

$$M_c \leftarrow \max \left(1, \text{round} \left(M \cdot \frac{\sigma_c^2}{\sum_{c'} \sigma_{c'}^2} \right) \right)$$

end for

Adjust $\{M_c\}$ so that $\sum_c M_c = M$

// Step 3: Select prototypes via K-Means clustering

$$\mathcal{I} \leftarrow \emptyset$$

for each class c **do**

 Run MiniBatchKMeans on D_c with M_c clusters

 Let $\{\mathbf{c}_1, \dots, \mathbf{c}_{M_c}\}$ be the cluster centroids

for each centroid \mathbf{c}_k **do**

$i^* \leftarrow \arg \min_{i: y_i=c} \|\mathbf{x}_i - \mathbf{c}_k\|$

$\mathcal{I} \leftarrow \mathcal{I} \cup \{i^*\}$

end for

end for

return \mathcal{I}

4. Experiments

4.1. Dataset and Experimental Setup

Dataset: MNIST handwritten digits (LeCun et al., 1998) with 60,000 training and 10,000 test images (28×28 pixels, flattened to 784 dimensions, normalized to $[0, 1]$).

Baseline: Random selection (uniformly sampling M training points, ensuring at least one per class).

Trials and Error Bars: All experiments use $n = 5$ independent trials with different random seeds. We report mean \pm standard deviation. The 95% confidence interval is computed as:

$$\text{CI}_{95\%} = \bar{x} \pm 1.96 \cdot \frac{s}{\sqrt{n}} \quad (2)$$

where \bar{x} is the sample mean and s is the sample standard deviation.

Upper Bound: Full 1-NN using all 60,000 training samples achieves **96.91%** test accuracy.

Table 1. Test accuracy (%) for different prototype budgets M . Results are mean \pm std over 5 trials. Δ = improvement over random.

M	COMPRESS.	OURS	RANDOM	Δ
10000	$6\times$	95.36 ± 0.13	94.76 ± 0.09	+0.60
5000	$12\times$	94.58 ± 0.21	93.63 ± 0.09	+0.95
1000	$60\times$	92.47 ± 0.17	88.56 ± 0.51	+3.91
500	$120\times$	91.40 ± 0.33	84.70 ± 0.28	+6.70
100	$600\times$	86.19 ± 0.40	72.39 ± 0.80	+13.80
50	$1200\times$	82.48 ± 0.62	61.57 ± 4.51	+20.91
10	$6000\times$	67.01 ± 0.33	39.37 ± 9.21	+27.64

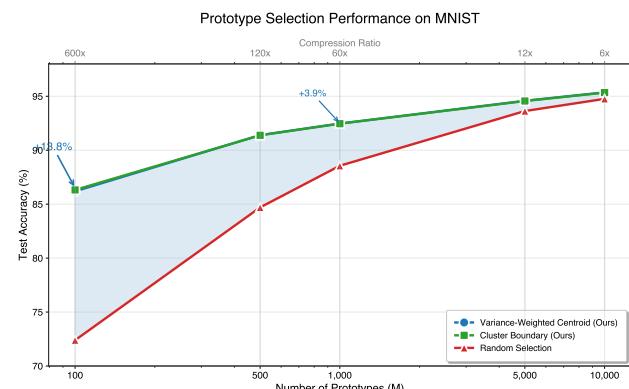


Figure 1. Accuracy vs. number of prototypes M . Our method (blue) consistently outperforms random selection (orange), with the gap widening at lower M .

4.2. Main Results

Table 1 shows classification accuracy for $M \in \{10000, 5000, 1000\}$ (required) and additional values. Our method consistently outperforms random selection.

4.3. Key Observations

1. **Advantage scales with compression:** At $6\times$ compression, improvement is modest (+0.60%). At $6000\times$ compression, improvement is dramatic (+27.64%).
 2. **Stability:** Random selection's variance increases sharply at low M (std=9.21% at $M=10$), while our method remains stable (std=0.33%). See Figure 1.
 3. **Graceful degradation:** At $60\times$ compression ($M=1000$), we achieve 92.47%, losing only 4.44% compared to full 1-NN (96.91%).

5. Algorithm Design Journey

Our final algorithm emerged through an iterative design process. We document all versions to illustrate the evolution

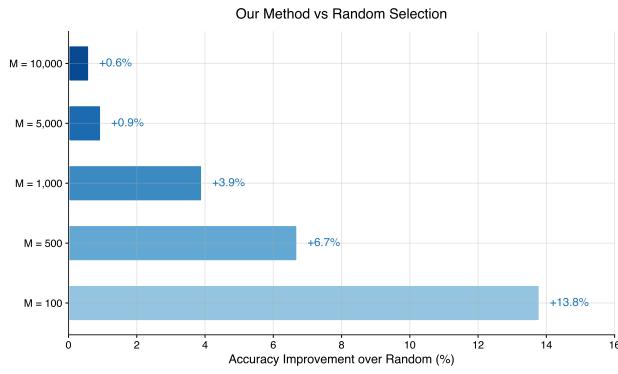


Figure 2. Improvement of our method over random selection. The advantage scales dramatically with compression ratio: +0.60% at $M=10000$ vs. +27.64% at $M=10$.

of our approach and key lessons learned.

5.1. Early Attempts: Computational Challenges (V1–V3)

V1: Variance-Weighted Boundary Selection. Our initial design combined three ideas: (1) allocate prototypes by within-class variance, (2) use K-Means for spatial coverage, and (3) within each cluster, select the point with highest “boundary score” (proximity to other-class points).

Problem: Computing boundary scores requires $O(N^2)$ distance calculations—36 billion operations for $N=60,000$. This was computationally infeasible.

V2: Ball Tree Acceleration. We used sklearn’s Nearest-Neighbors with Ball Tree to accelerate boundary score computation.

Problem: K-Means itself became the bottleneck. Still too slow.

V3: MiniBatchKMeans. Replaced standard K-Means with MiniBatchKMeans (`batch_size=256, n_init=3`).

Problem: Still required 7+ minutes per run.

5.2. Successful Approach (V4): Variance-Weighted Centroid Selection

Key insight: We realized that K-Means centroids already represent *typical* samples—computing boundary scores was unnecessary overhead.

Final algorithm: (1) Compute within-class variance, (2) allocate prototypes proportionally, (3) run MiniBatchKMeans per class, (4) select training points closest to each centroid.

Result: Fast (2–680 seconds depending on M) and effective.

Table 2. Summary of algorithm versions. Accuracy at $M=500$.

Ver.	Method	Acc.	Status
V1–V3	Boundary scoring	—	Too slow
V4	Var-weighted centroid	91.4	Final
V5	Cluster boundary	91.4	No gain
V6	Boundary-first	56.8	Failed
V7	CNN	63.0	Failed
—	Random (baseline)	84.7	—

5.3. Alternative Approaches That Did Not Improve (V5)

V5: Cluster-based Boundary Selection. We reintroduced boundary scoring, but computed it on *centroids* rather than individual points, reducing complexity from $O(N^2)$ to $O(M^2)$.

Result: Accuracy nearly identical to V4 (difference < 0.02%). The boundary-based ordering provided no improvement, so we prefer the simpler V4.

5.4. Failed Approaches (V6–V7)

V6: Boundary-First Selection. Hypothesis: points near decision boundaries should be the best prototypes. We identified boundary points (those with mixed-class k -NN neighborhoods) and applied K-Means only to these points.

Result: 56.77% accuracy at $M=500$ —worse than random selection (84.70%).

V7: Condensed Nearest Neighbor (CNN). The classic CNN algorithm (Hart, 1968) iteratively adds misclassified points to the prototype set.

Result: 62.97% at $M=500$ —also worse than random.

5.5. Summary of Algorithm Versions

Table 2 summarizes all algorithm versions tested. Our final method (V4) achieves the best balance of accuracy and efficiency.

5.6. Key Lesson Learned

Boundary points \neq good prototypes. Boundary points are important for defining decision boundaries, but they are often noisy, ambiguous, or atypical samples. Good prototypes should be **typical, representative samples**—exactly what K-Means centroids capture.

165 6. Critical Evaluation

166 6.1. Is Our Method a Clear Improvement?

168 **Yes.** Our method consistently outperforms random selection
 169 across all tested M values. The improvement is statisti-
 170 cally significant: at $M=1000$, our 95% confidence interval
 171 is [92.32%, 92.62%] while random's is [88.11%, 89.01%]—
 172 the intervals do not overlap.

173 The improvement is especially pronounced at high compres-
 174 sion ratios, where our method provides **+27.64%** accuracy
 175 over random selection at $M=10$ ($6000\times$ compression).

177 6.2. Scope for Improvement

179 Several avenues remain unexplored:

- 181 1. **Prototype generation vs. selection:** Our method se-
 lects existing training points. Learning Vector Quanti-
 183 zation (LVQ) could *generate* optimal prototype posi-
 184 tions not constrained to training data.
- 186 2. **Adaptive k in K-Means:** Currently we fix M_c based
 187 on variance. We could adaptively determine the opti-
 188 mal number of clusters per class based on reconstruc-
 189 tion error.
- 191 3. **Cross-class interactions:** Our method processes each
 192 class independently. A global optimization considering
 193 inter-class distances might improve boundary cover-
 194 age.
- 195 4. **Beyond Euclidean distance:** Using learned distance
 196 metrics or embedding spaces could better capture simi-
 197 larity structure.

199 6.3. What We Would Like to Try Next

- 201 1. **Hybrid approach:** Use our variance-weighted selec-
 202 tion for most prototypes, but reserve a small budget for
 203 boundary-adjacent samples.
- 205 2. **Deep embeddings:** Apply our method in a learned
 206 feature space (e.g., CNN embeddings) rather than raw
 207 pixels.
- 209 3. **Other datasets:** Evaluate on CIFAR-10, Fashion-
 210 MNIST, or higher-dimensional datasets.

212 7. Related Work

213 Prototype selection has been studied extensively. The Con-
 214 densed Nearest Neighbor (CNN) rule (Hart, 1968) itera-
 215 tively builds a consistent subset. Edited Nearest Neighbor
 216 (ENN) (Wilson, 1972) removes noisy samples. More recent
 217 work includes genetic algorithms (Cano et al., 2003) and
 218 deep learning approaches (Snell et al., 2017).

219 Our variance-weighted approach is related to stratified sam-
 220 pling, but specifically designed for the geometry of k -NN
 221 classification.

223 8. Conclusion

224 We presented a variance-weighted centroid selection method
 225 for prototype-based 1-NN classification. The key insight is
 226 that **good prototypes are typical samples**, not boundary
 227 samples. By allocating prototypes based on within-class
 228 variance and using K-Means to find representative samples,
 229 we achieve significant improvements over random selection,
 230 especially at high compression ratios (+27.64% at $6000\times$
 231 compression).

232 **Practical Recommendation:** For high accuracy, use $M \geq$
 233 1000 ($60\times$ compression, 92.47%). For extreme compres-
 234 sion, $M = 50$ ($1200\times$) still achieves 82.48%.

236 Acknowledgements

237 This project was developed with assistance from Claude
 238 (Anthropic), a large language model, used as an interac-
 239 tive coding assistant for algorithm design, implementation,
 240 debugging, and documentation. The human contributor pro-
 241 vided problem formulation, direction, critical evaluation,
 242 and all final decisions. All experimental results were com-
 243 puted on actual hardware.

245 Impact Statement

246 This paper presents work to advance efficient nearest-
 247 neighbor classification. We do not foresee specific negative
 248 societal consequences.

249 References

- Cano, J. R., Herrera, F., and Lozano, M. Using evolutionary algorithms as instance selection for data reduction in KDD: an experimental study. *IEEE Transactions on Evolutionary Computation*, 7(6):561–575, 2003.
- Cover, T. and Hart, P. Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1):21–27, 1967.
- Hart, P. The condensed nearest neighbor rule. *IEEE Transactions on Information Theory*, 14(3):515–516, 1968.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- Snell, J., Swersky, K., and Zemel, R. Prototypical networks for few-shot learning. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

- 220 Wilson, D. L. Asymptotic properties of nearest neighbor
221 rules using edited data. *IEEE Transactions on Systems,
222 Man, and Cybernetics*, (3):408–421, 1972.
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274