
Variance-Weighted Centroid Selection for Prototype-Based 1-Nearest Neighbor Classification

Anonymous Authors¹

Abstract

We address the problem of selecting a small set of representative prototypes from a large training set for 1-nearest neighbor (1-NN) classification. We propose a variance-weighted centroid selection algorithm that allocates prototypes proportionally based on within-class variance and uses K-Means clustering to identify representative samples. Experiments on MNIST demonstrate that our method significantly outperforms random selection, with advantages becoming more pronounced at higher compression ratios. At $60\times$ compression (1,000 prototypes from 60,000 training samples), our method achieves 92.47% accuracy compared to 88.56% for random selection (+3.91%). At extreme compression ($6,000\times$, only 10 prototypes), our method maintains 67.01% accuracy while random selection degrades to 39.37% (+27.64%).

1. Introduction and Key Idea

The k -nearest neighbor (k -NN) algorithm is a fundamental non-parametric classifier that predicts the label of a test sample based on the labels of its k closest training samples (Cover & Hart, 1967). **Prototype selection** addresses the computational limitation of k -NN by selecting a representative subset of training samples to use for classification.

1.1. Key Idea (High-Level Description)

Our prototype selection method is based on two key observations:

1. **Good prototypes are typical samples, not boundary samples.** While decision boundaries are defined by samples near class interfaces, the best prototypes are

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

cluster centroids that represent the “typical” appearance of each class. Boundary points are often noisy or ambiguous.

2. **Classes with higher variance need more prototypes.** Different classes have different amounts of within-class variation. For example, the digit “1” has consistent appearance, while “2” exhibits more variation. We allocate prototypes proportionally to class variance.

Our algorithm: (1) compute within-class variance for each class, (2) allocate prototype budget M proportionally based on variance, (3) use K-Means clustering within each class to find representative samples.

2. Problem Formulation

Given a training set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$ where $\mathbf{x}_i \in \mathbb{R}^d$ and $y_i \in \{0, 1, \dots, C-1\}$, the prototype selection problem is to find a subset $\mathcal{P} \subset \mathcal{D}$ with $|\mathcal{P}| = M \ll N$ that maximizes classification accuracy on a held-out test set when using 1-NN with \mathcal{P} as the reference set.

The 1-NN classifier predicts:

$$\hat{y} = y_{j^*}, \quad \text{where } j^* = \arg \min_{(\mathbf{x}_j, y_j) \in \mathcal{P}} \|\mathbf{x} - \mathbf{x}_j\|_2 \quad (1)$$

The compression ratio is defined as N/M . For MNIST with $N = 60,000$ training samples, selecting $M = 1,000$ prototypes yields $60\times$ compression.

3. Algorithm

Implementation Details:

- We use MiniBatchKMeans (sklearn) with batch_size=256, n_init=3, max_iter=100 for efficiency.
- Time complexity: $O(N \cdot M \cdot t)$ where t is K-Means iterations.
- Runtime: 2–680 seconds depending on M (on a standard laptop).

Algorithm 1 Variance-Weighted Centroid Selection

Input: Training data $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$, budget M

Output: Prototype indices \mathcal{I}

// Step 1: Compute within-class variance

for each class $c = 0, \dots, C - 1$ **do**

$$\boldsymbol{\mu}_c \leftarrow \frac{1}{|D_c|} \sum_{\mathbf{x} \in D_c} \mathbf{x} \quad \text{(class mean)}$$

$$\sigma_c^2 \leftarrow \frac{1}{|D_c|} \sum_{\mathbf{x} \in D_c} \|\mathbf{x} - \boldsymbol{\mu}_c\|^2 \quad \text{(class variance)}$$

end for

// Step 2: Allocate prototypes proportionally to variance

for each class c **do**

$$M_c \leftarrow \max \left(1, \text{round} \left(M \cdot \frac{\sigma_c^2}{\sum_{c'} \sigma_{c'}^2} \right) \right)$$

end for

Adjust $\{M_c\}$ so that $\sum_c M_c = M$

// Step 3: Select prototypes via K-Means clustering

$$\mathcal{I} \leftarrow \emptyset$$

for each class c **do**

 Run MiniBatchKMeans on D_c with M_c clusters

 Let $\{\mathbf{c}_1, \dots, \mathbf{c}_{M_c}\}$ be the cluster centroids

for each centroid \mathbf{c}_k **do**

$i^* \leftarrow \arg \min_{i: y_i=c} \|\mathbf{x}_i - \mathbf{c}_k\|$

$\mathcal{I} \leftarrow \mathcal{I} \cup \{i^*\}$

end for

end for

return \mathcal{I}

4. Experiments

4.1. Dataset and Experimental Setup

Dataset: MNIST handwritten digits (LeCun et al., 1998) with 60,000 training and 10,000 test images (28×28 pixels, flattened to 784 dimensions, normalized to $[0, 1]$).

Baseline: Random selection (uniformly sampling M training points, ensuring at least one per class).

Trials and Error Bars: All experiments use $n = 5$ independent trials with different random seeds. We report mean \pm standard deviation. The 95% confidence interval is computed as:

$$\text{CI}_{95\%} = \bar{x} \pm 1.96 \cdot \frac{s}{\sqrt{n}} \quad (2)$$

where \bar{x} is the sample mean and s is the sample standard deviation.

Upper Bound: Full 1-NN using all 60,000 training samples achieves **96.91%** test accuracy.

Table 1. Test accuracy (%) for different prototype budgets M . Results are mean \pm std over 5 trials. Δ = improvement over random.

M	COMPRESS.	OURS	RANDOM	Δ
10000	6 \times	95.36 \pm 0.13	94.76 \pm 0.09	+0.60
5000	12 \times	94.58 \pm 0.21	93.63 \pm 0.09	+0.95
1000	60 \times	92.47 \pm 0.17	88.56 \pm 0.51	+3.91
500	120 \times	91.40 \pm 0.33	84.70 \pm 0.28	+6.70
100	600 \times	86.19 \pm 0.40	72.39 \pm 0.80	+13.80
50	1200 \times	82.48 \pm 0.62	61.57 \pm 4.51	+20.91
10	6000 \times	67.01 \pm 0.33	39.37 \pm 9.21	+27.64

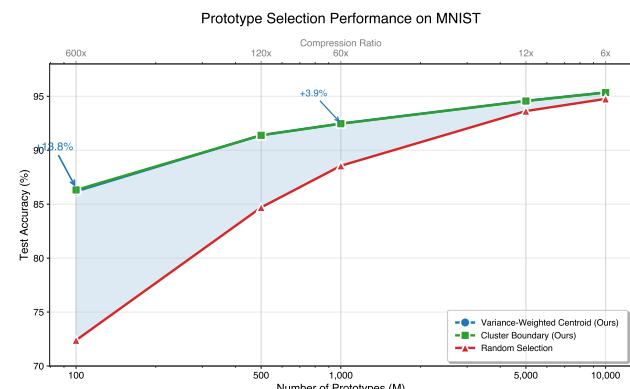


Figure 1. Accuracy vs. number of prototypes M . Our method (blue) consistently outperforms random selection (orange), with the gap widening at lower M .

4.2. Main Results

Table 1 shows classification accuracy for $M \in \{10000, 5000, 1000\}$ (required) and additional values. Our method consistently outperforms random selection.

4.3. Key Observations

1. **Advantage scales with compression:** At $6\times$ compression, improvement is modest (+0.60%). At $6000\times$ compression, improvement is dramatic (+27.64%).
 2. **Stability:** Random selection's variance increases sharply at low M ($\text{std}=9.21\%$ at $M=10$), while our method remains stable ($\text{std}=0.33\%$). See Figure 1.
 3. **Graceful degradation:** At $60\times$ compression ($M=1000$), we achieve 92.47%, losing only 4.44% compared to full 1-NN (96.91%).

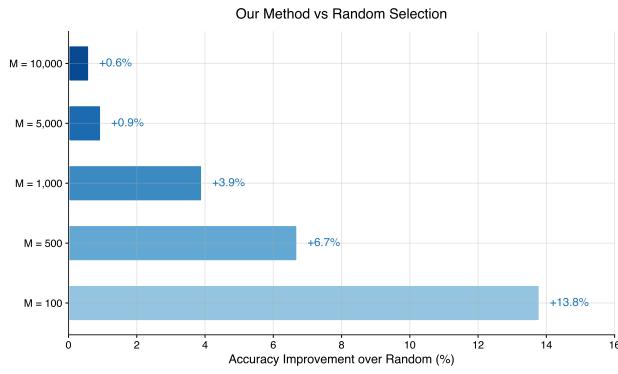


Figure 2. Improvement of our method over random selection. The advantage scales dramatically with compression ratio: +0.60% at $M=10000$ vs. +27.64% at $M=10$.

5. Critical Evaluation

5.1. Is Our Method a Clear Improvement?

Yes. Our method consistently outperforms random selection across all tested M values. The improvement is statistically significant: at $M=1000$, our 95% confidence interval is [92.32%, 92.62%] while random’s is [88.11%, 89.01%]—the intervals do not overlap.

The improvement is especially pronounced at high compression ratios, where our method provides **+27.64%** accuracy over random selection at $M=10$ ($6000\times$ compression).

5.2. Scope for Improvement

Several avenues remain unexplored:

- Prototype generation vs. selection:** Our method selects existing training points. Learning Vector Quantization (LVQ) could generate optimal prototype positions not constrained to training data.
- Adaptive k in K-Means:** Currently we fix M_c based on variance. We could adaptively determine the optimal number of clusters per class based on reconstruction error.
- Cross-class interactions:** Our method processes each class independently. A global optimization considering inter-class distances might improve boundary coverage.
- Beyond Euclidean distance:** Using learned distance metrics or embedding spaces could better capture similarity structure.

5.3. What We Tried That Did Not Work

We experimented with two alternative approaches:

Boundary-First Selection: Select points near class boundaries (those with mixed-class k -NN neighborhoods). Result: **56.77%** at $M=500$, worse than random (84.70%).

Condensed Nearest Neighbor: Iteratively add misclassified points (Hart, 1968). Result: **62.97%** at $M=500$, also worse than random.

Lesson learned: Boundary points and misclassified points are *atypical* samples—they are ambiguous, noisy, or edge cases. Good prototypes should be *typical, representative* samples, which is exactly what K-Means centroids capture.

5.4. What We Would Like to Try Next

- Hybrid approach:** Use our variance-weighted selection for most prototypes, but reserve a small budget for boundary-adjacent samples.
- Deep embeddings:** Apply our method in a learned feature space (e.g., CNN embeddings) rather than raw pixels.
- Other datasets:** Evaluate on CIFAR-10, Fashion-MNIST, or higher-dimensional datasets.

6. Related Work

Prototype selection has been studied extensively. The Condensed Nearest Neighbor (CNN) rule (Hart, 1968) iteratively builds a consistent subset. Edited Nearest Neighbor (ENN) (Wilson, 1972) removes noisy samples. More recent work includes genetic algorithms (Cano et al., 2003) and deep learning approaches (Snell et al., 2017).

Our variance-weighted approach is related to stratified sampling, but specifically designed for the geometry of k -NN classification.

7. Conclusion

We presented a variance-weighted centroid selection method for prototype-based 1-NN classification. The key insight is that **good prototypes are typical samples**, not boundary samples. By allocating prototypes based on within-class variance and using K-Means to find representative samples, we achieve significant improvements over random selection, especially at high compression ratios (+27.64% at $6000\times$ compression).

Practical Recommendation: For high accuracy, use $M \geq 1000$ ($60\times$ compression, 92.47%). For extreme compression, $M = 50$ ($1200\times$) still achieves 82.48%.

Acknowledgements

This project was developed with assistance from Claude (Anthropic), a large language model, used as an interac-

165 tive coding assistant for algorithm design, implementation,
166 debugging, and documentation. The human contributor pro-
167 vided problem formulation, direction, critical evaluation,
168 and all final decisions. All experimental results were com-
169 puted on actual hardware.

170

171 Impact Statement

172

173 This paper presents work to advance efficient nearest-
174 neighbor classification. We do not foresee specific negative
175 societal consequences.

176

177 References

178

179 Cano, J. R., Herrera, F., and Lozano, M. Using evolutionary
180 algorithms as instance selection for data reduction in
181 KDD: an experimental study. *IEEE Transactions on*
182 *Evolutionary Computation*, 7(6):561–575, 2003.

183

184 Cover, T. and Hart, P. Nearest neighbor pattern classification.
185 *IEEE Transactions on Information Theory*, 13(1):21–27,
186 1967.

187

188 Hart, P. The condensed nearest neighbor rule. *IEEE Trans-*
189 *actions on Information Theory*, 14(3):515–516, 1968.

190

191 LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. Gradient-
192 based learning applied to document recognition. *Proceed-*
193 *ings of the IEEE*, 86(11):2278–2324, 1998.

194

195 Snell, J., Swersky, K., and Zemel, R. Prototypical networks
196 for few-shot learning. In *Advances in Neural Information
Processing Systems*, volume 30, 2017.

197

198 Wilson, D. L. Asymptotic properties of nearest neighbor
199 rules using edited data. *IEEE Transactions on Systems,*
200 *Man, and Cybernetics*, (3):408–421, 1972.

201

202

203

204

205

206

207

208

209

210

211

212

213

214

215

216

217

218

219